

# Capstone Project

## Face Emotion Recognition

**Name- Pranav Singhal**  
**Batch-AlmaBetter Pro**

# Problem Description

The Indian education landscape has been undergoing rapid changes for the past 10 years owing to the advancement of web-based learning services, specifically, eLearning platforms.

In a physical classroom during a lecturing teacher can see the faces and assess the emotion of the class and tune their lecture accordingly, whether he is going fast or slow. He can identify students who need special attention. Digital classrooms are conducted via video telephony software program (exZoom) where it's not possible for medium scale class (25-50) to see all students and access the mood. Because of this drawback, students are not focusing on content due to lack of surveillance




# How can Face Emotion Recognition with Deep Learning help? AI

While digital platforms have limitations in terms of physical surveillance but it comes with the power of data and machines which can work for you. It provides data in the form of video, audio, and texts which can be analysed using deep learning algorithms. Deep learning backed system not only solves the surveillance issue, but it also removes the human bias from the system, and all information is no longer in the teacher's brain rather translated in numbers that can be analyzed and tracked.



# Data – FER2013

Kaggle Link- <https://www.kaggle.com/msambare/fer2013>



The data consists of 48x48 pixel grayscale images of faces. The faces have been automatically registered so that the face is more or less centered and occupies about the same amount of space in each image.

The task is to categorize each face based on the emotion shown in the facial expression into one of seven categories (0=Angry, 1=Disgust, 2=Fear, 3=Happy, 4=Sad, 5=Surprise, 6=Neutral). The training set consists of 28,709 examples and the public test set consists of 3,589 examples.








#### Data Explorer

56.51 MB

- test
  - angry
  - disgust
  - fear
  - happy
  - neutral
  - sad
  - surprise
- train
  - angry
  - disgust
  - fear
  - happy
  - neutral
  - sad
  - surprise

#### < test (7 directories)



 angry 958 files	 disgust 111 files	 fear 1024 files	 happy 1774 files	 neutral 1233 files
 sad 1247 files	 surprise 831 files			

Data Structure : I would be creating a validation Data set apart from these test and train sets

# Loading Data in Notebook

```
training_gen=ImageDataGenerator(rescale=1./255)                                     # Creating Image generators, for all train, validation, and test set
validation_gen=ImageDataGenerator(rescale=1./255)
testing_gen=ImageDataGenerator(rescale=1./255)

train_gen_3ch=training_gen.flow_from_directory('archive/train',                      # Creating Training Dataset with 3 channels
                                                target_size=(48,48),
                                                batch_size=32,
                                                class_mode='categorical')
train_gen=training_gen.flow_from_directory('archive/train',                        # Creating Training Dataset
                                           target_size=(48,48),
                                           batch_size=32,
                                           color_mode='grayscale',
                                           class_mode='categorical')
valid_gen=validation_gen.flow_from_directory('archive/validation',                # Creating Validation Set
                                              target_size=(48,48),
                                              batch_size=32,
                                              color_mode='grayscale',
                                              class_mode='categorical')
test_gen=testing_gen.flow_from_directory('archive/test',                          # Creating Test Set
                                         target_size=(48,48),
                                         batch_size=32,
                                         color_mode='grayscale',
                                         class_mode='categorical',
                                         shuffle= False)

Found 28207 images belonging to 7 classes.
Found 28207 images belonging to 7 classes.
Found 7178 images belonging to 7 classes.
Found 502 images belonging to 7 classes.
```

Loading Data and Splitting it into train, test and Validation set and assigning Class to each of them on the basis of folder they are present in, here I am setting the target size as 48 X 48 and keeping only one channel out of 3 by providing color mode as Grayscale

# Plotting Data

```
[ ] imgs,labels=next(train_gen_3ch)
list_of_keys=['angry', 'disgust', 'fear', 'happy', 'neutral', 'sad', 'surprise']
def plotImages(images_arr):
    fig, axes = plt.subplots(1,32,figsize=(32,2))
    axes = axes.flatten()
    for img, ax in zip(images_arr,axes):
        ax.imshow(img)
        ax.axis('off')
    plt.tight_layout()
    plt.show()
plotImages(imgs)
for i in labels[:32]:
    for num, j in enumerate(i):
        if j==1:
            print(list_of_keys[num], end='  ')
num+=1
```

```
# Extracting next batch to plot it, this batch of photos are selected randomly
# list of Classes

# Printing the labels below the respective image

# Labelling seems to be spot on
```



Here I am Plotting one Batch of train set with their respective class labels, and by looking at the Images and their labels, I can conclude that Labeling is really very accurate.

# CNN Model

```

cnn_model= Sequential([                                     # Usi
    ng Keras Sequential model to build out CNN
    Conv2D(32, kernel_size=(3,3), activation='relu',input_shape=(48,48,1)),      # Fir
    st Layer of convolutional Nural network, input shape is provided only at this instance.
    Conv2D(64, kernel_size=(3,3), activation='relu'),
    MaxPool2D(pool_size=(2,2)),                                                  # Max
    pool Layer
    Dropout(0.05),

    Conv2D(128, kernel_size=(3,3), activation='relu'),
    MaxPool2D(pool_size=(2,2)),
    Dropout(0.05),

    Conv2D(256, kernel_size=(3,3), activation='relu'),
    MaxPool2D(pool_size=(2,2)),
    Dropout(0.05),

    Conv2D(512, kernel_size=(3,3), activation='relu'),
    MaxPool2D(pool_size=(2,2)),
    Dropout(0.05),

    Flatten(),                                                                    # Fla
    ttening out all the layers
    Dense(1024, activation='relu'),
    Dense(7, activation='softmax')]
)
cnn_model.summary()

```



# Model Summary

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 46, 46, 32)	320
conv2d_1 (Conv2D)	(None, 44, 44, 64)	18496
max_pooling2d (MaxPooling2D)	(None, 22, 22, 64)	0
dropout (Dropout)	(None, 22, 22, 64)	0
conv2d_2 (Conv2D)	(None, 20, 20, 128)	73856
max_pooling2d_1 (MaxPooling2D)	(None, 10, 10, 128)	0
dropout_1 (Dropout)	(None, 10, 10, 128)	0
conv2d_3 (Conv2D)	(None, 8, 8, 256)	295168
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 256)	0
dropout_2 (Dropout)	(None, 4, 4, 256)	0
conv2d_4 (Conv2D)	(None, 2, 2, 512)	1180160
max_pooling2d_3 (MaxPooling2D)	(None, 1, 1, 512)	0
dropout_3 (Dropout)	(None, 1, 1, 512)	0
flatten (Flatten)	(None, 512)	0
dense_3 (Dense)	(None, 1024)	525312
dense_4 (Dense)	(None, 7)	7175
Total params: 2,100,487		
Trainable params: 2,100,487		
Non-trainable params: 0		

# Compiling and Fitting the Model

```
checkpoint = ModelCheckpoint('./my_model.h5', monitor='val_acc', verbose=1, save_best_only=True, mode='max')

early_stopping=EarlyStopping(monitor='val_loss',
                              min_delta=0.1,
                              patience=4,
                              verbose=2,
                              restore_best_weights=True)

decay_lr= ReduceLROnPlateau(monitor='val_loss',
                             factor=0.2,
                             patience=3,
                             verbose=1,
                             min_delta=0.0001)

callbacks=[early_stopping,checkpoint,decay_lr]
```

```
cnn_model.compile(loss='categorical_crossentropy',optimizer='adam', metrics=['accuracy'])
piling the Model
cnn_model.fit(train_gen, epochs=20, verbose=2,validation_data=valid_gen, callbacks=callbacks)
```

# Com

Here I have used Adam Optimizer, and Loss to be Categorical cross entropy

For training the model to train dataset, a training and validation set is provided and callbacks like early stopping and decay learning rate are added for better training

# Accuracy on train and Validation Set

```

Epoch 1/20
882/882 - 186s - loss: 1.7513 - accuracy: 0.2793 - val_loss: 1.5645 - val_accuracy: 0.3835
WARNING:tensorflow:Can save best model only with val_acc available, skipping.
Epoch 2/20
882/882 - 131s - loss: 1.4599 - accuracy: 0.4289 - val_loss: 1.3645 - val_accuracy: 0.4710
WARNING:tensorflow:Can save best model only with val_acc available, skipping.
Epoch 3/20
882/882 - 130s - loss: 1.2926 - accuracy: 0.5008 - val_loss: 1.2440 - val_accuracy: 0.5228
WARNING:tensorflow:Can save best model only with val_acc available, skipping.
Epoch 4/20
882/882 - 130s - loss: 1.1842 - accuracy: 0.5460 - val_loss: 1.1845 - val_accuracy: 0.5460
WARNING:tensorflow:Can save best model only with val_acc available, skipping.
Epoch 5/20
882/882 - 130s - loss: 1.1053 - accuracy: 0.5818 - val_loss: 1.1593 - val_accuracy: 0.5538
WARNING:tensorflow:Can save best model only with val_acc available, skipping.
Epoch 6/20
882/882 - 131s - loss: 1.0353 - accuracy: 0.6094 - val_loss: 1.1470 - val_accuracy: 0.5596
WARNING:tensorflow:Can save best model only with val_acc available, skipping.
Epoch 7/20
882/882 - 131s - loss: 0.9697 - accuracy: 0.6333 - val_loss: 1.1301 - val_accuracy: 0.5747
WARNING:tensorflow:Can save best model only with val_acc available, skipping.
Epoch 8/20
882/882 - 131s - loss: 0.9109 - accuracy: 0.6581 - val_loss: 1.1307 - val_accuracy: 0.5819
WARNING:tensorflow:Can save best model only with val_acc available, skipping.
Epoch 9/20
882/882 - 130s - loss: 0.8440 - accuracy: 0.6821 - val_loss: 1.1769 - val_accuracy: 0.5804
WARNING:tensorflow:Can save best model only with val_acc available, skipping.
Epoch 10/20
882/882 - 130s - loss: 0.7884 - accuracy: 0.7049 - val_loss: 1.1951 - val_accuracy: 0.5761
WARNING:tensorflow:Can save best model only with val_acc available, skipping.

Epoch 00010: ReduceLROnPlateau reducing learning rate to 0.00020000000949949026.
Epoch 11/20
882/882 - 130s - loss: 0.6017 - accuracy: 0.7768 - val_loss: 1.2628 - val_accuracy: 0.5921
Restoring model weights from the end of the best epoch.
WARNING:tensorflow:Can save best model only with val_acc available, skipping.
Epoch 00011: early stopping

```

- Maximum accuracy for training set is 70% while for validation set is only 60%
- Thus the model is under fitting.

# Predicting Facial expression

```
predictions=cnn_model.predict(test_gen)
from sklearn.metrics import confusion_matrix, classification_report
cm=confusion_matrix(y_pred=np.argmax(predictions, axis=-1), y_true=test_gen.classes)
cm
```

```
array([[42,  1,  5,  9,  8,  8,  2],
       [30,  8, 11, 10,  1, 13,  2],
       [ 6,  0, 10,  2,  2,  8,  4],
       [ 0,  0,  3, 67,  2,  7,  1],
       [10,  0, 10, 13, 23, 20,  4],
       [10,  1,  6,  7, 11, 44,  1],
       [ 0,  0,  6,  6,  1,  2, 65]], dtype=int64)
```

```
from sklearn.metrics import accuracy_score
aoc=accuracy_score(y_pred=np.argmax(predictions, axis=-1),y_true=test_gen.classes)
aoc                                     # I am able to get an accuracy of only 50% Lets try with a pretrained model
```

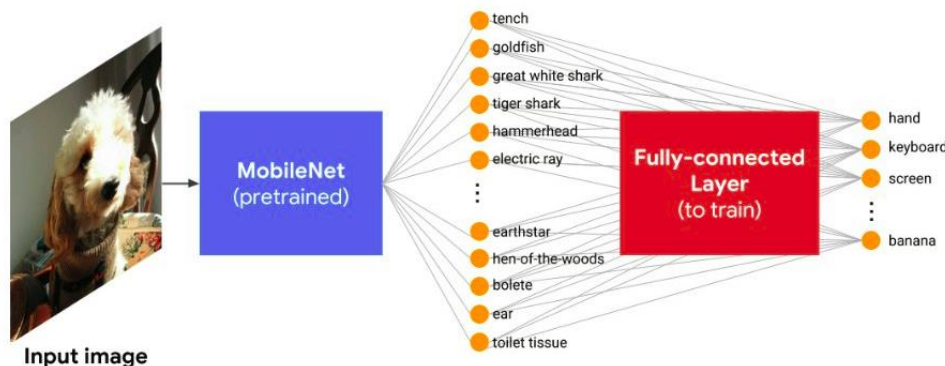
```
0.5159362549800797
```

```
cnn_model.save('my_model.h5')
```

Here on Predicting the Facial Expression of test data set it can be seen that Model is showing average performance as Accuracy score is only 51%

# Transfer Learning

Transfer learning, used in machine learning, is the reuse of a pre-trained model on a new problem. In transfer learning, a machine exploits the knowledge gained from a previous task to improve generalization about another. For example, in training a classifier to predict whether an image contains food, you could use the knowledge it gained during training to recognize drinks.



Architecture of our image recognition model. We trained a fully-connected layer added to the pretrained MobileNet.

# Loading Mobile Net Model and Making a few amendments to it

```
import random
random.shuffle(train_set)
```

```
X=[]
y=[]
for features, labels in train_set:
    X.append(features)
    y.append(labels)
X=np.array(X).reshape(-1,224,224,3)
```

```
from sklearn.model_selection import train_test_split
X_train, X_validation, y_train, y_validation= train_test_split(X,y,test_size=0.1)
```

```
final_output=layers.Dense(128)(base_output)
final_output=layers.Activation('relu')(final_output)
final_output=layers.Dense(128)(final_output)
final_output=layers.Activation('relu')(final_output)
final_output=layers.Dense(7, activation='softmax')(final_output)
```

```
final_output
```

```
<KerasTensor: shape=(None, 7) dtype=float32 (created by layer 'dense_2')>
```

```
final_model=keras.Model(inputs=base_input, outputs=final_output)
```

```
for layer in final_model.layers[:-23]:
    layer.trainable = False
```

Creating Train and test set yet again as Mobile Net architecture takes input of (None,224,224,3), thus keeping all 3 channels

Making Few amendments to the final layer of model, instead on classifying 1000 objects, now it will only be able to classify 7.

All I am freezing few weights and making them un-trainable

Conv_1 (Conv2D)	(None, 7, 7, 1280)	409600	block_16_project_BN[0][0]
Conv_1_bn (BatchNormalization)	(None, 7, 7, 1280)	5120	Conv_1[0][0]
out_relu (ReLU)	(None, 7, 7, 1280)	0	Conv_1_bn[0][0]
global_average_pooling2d (GlobalAveragePooling2D)	(None, 1280)	0	out_relu[0][0]
dense (Dense)	(None, 128)	163968	global_average_pooling2d[0][0]
activation (Activation)	(None, 128)	0	dense[0][0]
dense_1 (Dense)	(None, 128)	16512	activation[0][0]
activation_1 (Activation)	(None, 128)	0	dense_1[0][0]
dense_2 (Dense)	(None, 7)	903	activation_1[0][0]
=====			
Total params: 2,439,367			
Trainable params: 1,231,943			
Non-trainable params: 1,207,424			

- Model Summary : As Few of the Parameters are at freeze, there are only some which are trainable

# Model Performance

```
final_model.compile(loss='categorical_crossentropy', optimizer=Adam(lr=0.0001), metrics=['accuracy'])
final_model.fit(X_train, y_train, epochs=20, verbose=2, validation_data=(X_validation, y_validation))
```

```
Epoch 1/20
794/794 - 406s - loss: 1.5786 - accuracy: 0.3765 - val_loss: 2.2255 - val_accuracy: 0.1581
Epoch 2/20
794/794 - 401s - loss: 1.3411 - accuracy: 0.4864 - val_loss: 2.0981 - val_accuracy: 0.2233
Epoch 3/20
794/794 - 401s - loss: 1.1818 - accuracy: 0.5559 - val_loss: 2.3961 - val_accuracy: 0.2205
Epoch 4/20
794/794 - 401s - loss: 1.0205 - accuracy: 0.6248 - val_loss: 1.7945 - val_accuracy: 0.3559
Epoch 5/20
794/794 - 403s - loss: 0.8426 - accuracy: 0.6978 - val_loss: 1.5884 - val_accuracy: 0.4325
Epoch 6/20
794/794 - 406s - loss: 0.6698 - accuracy: 0.7661 - val_loss: 1.8180 - val_accuracy: 0.4403
Epoch 7/20
794/794 - 400s - loss: 0.5158 - accuracy: 0.8242 - val_loss: 1.9420 - val_accuracy: 0.4172
Epoch 8/20
794/794 - 396s - loss: 0.3797 - accuracy: 0.8767 - val_loss: 2.0276 - val_accuracy: 0.4537
Epoch 9/20
794/794 - 396s - loss: 0.2783 - accuracy: 0.9123 - val_loss: 2.0764 - val_accuracy: 0.4782
Epoch 10/20
794/794 - 394s - loss: 0.2056 - accuracy: 0.9379 - val_loss: 2.5549 - val_accuracy: 0.4488
Epoch 11/20
794/794 - 395s - loss: 0.1700 - accuracy: 0.9480 - val_loss: 2.7138 - val_accuracy: 0.4619
Epoch 12/20
794/794 - 393s - loss: 0.1461 - accuracy: 0.9560 - val_loss: 2.6793 - val_accuracy: 0.4697
Epoch 13/20
794/794 - 395s - loss: 0.1244 - accuracy: 0.9619 - val_loss: 3.2451 - val_accuracy: 0.4169
Epoch 14/20
794/794 - 394s - loss: 0.1162 - accuracy: 0.9641 - val_loss: 2.9203 - val_accuracy: 0.4413
Epoch 15/20
794/794 - 395s - loss: 0.1090 - accuracy: 0.9657 - val_loss: 3.0480 - val_accuracy: 0.4307
Epoch 16/20
794/794 - 394s - loss: 0.1055 - accuracy: 0.9672 - val_loss: 3.0312 - val_accuracy: 0.4541
Epoch 17/20
794/794 - 394s - loss: 0.0931 - accuracy: 0.9721 - val_loss: 3.1519 - val_accuracy: 0.4225
Epoch 18/20
794/794 - 395s - loss: 0.0964 - accuracy: 0.9688 - val_loss: 3.1034 - val_accuracy: 0.4378
Epoch 19/20
794/794 - 394s - loss: 0.0972 - accuracy: 0.9681 - val_loss: 3.3477 - val_accuracy: 0.4495
Epoch 20/20
794/794 - 408s - loss: 0.0837 - accuracy: 0.9733 - val_loss: 3.5165 - val_accuracy: 0.4413
<keras.callbacks.History at 0x22e1d4f9580>
```

This Time Around the Model is overfitting as Training Accuracy is around 97.3 % while validation accuracy is only 44 % thus this model is not Robust enough and Model I made previously was better



# Creating Web App Using Streamlit

Streamlit is an open-source python framework for building web apps for Machine Learning and Data Science. We can instantly develop web apps and deploy them easily using Streamlit. Streamlit allows you to write an app the same way you write a python code. Streamlit makes it seamless to work on the interactive loop of coding and viewing results in the web app.




# Deploy Model Using AWS

Cloud computing is the on-demand delivery of IT resources over the Internet with pay-as-you-go pricing. Instead of buying, owning, and maintaining physical data centers and servers, you can access technology services, such as computing power, storage, and databases, on an as-needed basis from a cloud provider like Amazon Web Services (AWS).



Amazon Web Services (AWS) is the world's most comprehensive and broadly adopted cloud platform, offering over 200 fully featured services from data centers globally. Millions of customers—including the fastest-growing startups, largest enterprises, and leading government agencies—are using AWS.







## Conclusion

In the End we have a CNN Model which can detect the facial expressions, using which we can monitor the facial expression of all the students in the class simultaneously, using this feed teacher can assess the emotion of the class and tune their lecture accordingly, whether he is going fast or slow. He can identify students who need special attention.

THANK YOU